

WHEN PERSONALIZATION DETERS DELEGATION TO ALGORITHMS: EVIDENCE FROM A PREREGISTERED EXPERIMENT*

Lara Suraci

DG Cities

Irenaeus Wolff†

*Thurgau Institute of Economics
(TWI) / University of Konstanz*

This version: 14th April, 2026

Abstract

Algorithmic decision aids increasingly act on users' behalf, and in many instances users can either retain control or delegate decisions to an algorithm altogether. We test whether preference-based personalization increases or decreases willingness to delegate compared with an objective optimization criterion. In a preregistered between-subjects experiment ($N = 233$), participants repeatedly chose whether to decide themselves or to delegate to an algorithm. In one treatment, the delegation option is an optimization-based algorithm selecting the option with the highest expected value. In the other treatment, it is a preference-based (personalized) algorithm selecting according to participants' previously elicited preferences. We find no evidence that preference-based personalization increases delegation relative to expected-value optimization throughout the task. Instead, personalization deters delegation in two ways. First, at first contact participants are significantly less likely to delegate to the preference-based algorithm than to the optimization-based algorithm; this difference is confined to the very first decision and disappears immediately thereafter, with delegation rates statistically indistinguishable across algorithm types in subsequent decisions. Second, categorical non-delegation (never delegating) is substantially more prevalent under preference-based personalization (15.7%) than under objective optimization (3.4%). All in all, our results are consistent with the idea that preference-matching can raise initial hesitation and heighten agency-related concerns for a subset of users, functioning as an acceptance barrier even when personalization is normatively appealing.

*We are grateful to all the many people who helped us write this paper, first and foremost Chris Starmer and Robin Cubitt, to whom we owe several aspects of the design, as well as the lively research group at the TWI and participants of numerous workshops and conferences.

†Corresponding author. Thurgau Institute of Economics, Hafenstrasse 6, 8280 Kreuzlingen, Switzerland, wolff@twi-kreuzlingen.ch.

Keywords: AI delegation, human-AI delegation, personalization, algorithm appreciation, algorithm aversion, trust in AI, financial decision-making, preregistered experiment.

1 Introduction

Algorithmic systems increasingly make and implement decisions on users’ behalf in both everyday and high-stakes contexts, from automated driving to autonomous trading and smart-home control. Public perceptions of algorithmic decision-making vary substantially across domains and are shaped by perceived fairness, accountability, privacy, trust, and usefulness (Aysolmaz et al., 2023). In many such settings, users face a concrete choice: they can retain control and decide themselves, or delegate the decision to an algorithm. Delegation can reduce cognitive effort and simplify complex choice tasks, but it can also reshape perceived agency and responsibility for outcomes. Recent literature shows that people do delegate consequential choices to AI systems, and that this delegation depends on the decision context (Candrian and Scherer, 2022). Note that throughout the paper, we refer to these systems interchangeably as algorithms, algorithmic advice, and algorithmic decision aids. This reflects the practical reality that, from a user’s perspective, the boundary between AI and algorithmic decision rules is often opaque.

A common assumption is that *personalization* should increase acceptance. If individuals differ in preferences, then a decision aid that incorporates those preferences should, in principle, deliver higher subjective utility than a one-size-fits-all rule. In financial decision support, for example, systems often elicit risk preferences and translate them into portfolio recommendations. The key empirical question, however, is not whether personalization sounds attractive in principle, but whether it actually increases willingness to delegate.

Existing evidence leaves that question open. In robo-advisor research, much of the evidence focuses on adoption correlates and advisor-use patterns (for example, Mahmud et al., 2025), rather than on direct contrasts between preference-based personalization and objective delegation rules. Adjacent work has shown that giving users even slight control over an algorithm’s output can increase adoption (Dietvorst et al., 2018), and that conversational onboarding and autonomy-preserving designs shape trust and uptake (Hildebrand and Bergner, 2021; Fink et al., 2024). However, these studies vary the interface, degree of user control, or communication format—not the nature of the underlying decision rule itself. To our knowledge, no study has experimentally tested whether making an algorithm’s decision rule preference-based, rather than objectively optimizing, changes actual delegation behavior. More generally, behavioral evidence shows that delegation

to algorithms is shaped by more than expected performance: people sometimes avoid algorithms and sometimes prefer them based on the context (Dietvorst et al., 2018; Logg et al., 2019); they need to understand how recommendations are generated before accepting them (Yeomans et al., 2019); and acceptance depends on moderators such as perceived control and task features (Mahmud et al., 2022; Chugunova and Sele, 2022).

Preference-based personalization can therefore be double-edged. While it may increase perceived fit, it may also encroach on users’ sense of choice authority: a system that claims to decide *as you would* substitutes for personal judgment more directly than an impersonal optimization rule, making delegation feel like ceding active control rather than outsourcing a computation. Our own post-task evidence is consistent with this interpretation: many categorical non-delegators explicitly report wanting to *keep* responsibility for the choice (but not the outcome), suggesting that the barrier is not fear of being held responsible but rather a positive valuation of exercising choice authority. This reading aligns with recent findings that transparency can backfire by triggering negative social interpretations (Chen et al., 2024), that transparency can induce resistance in algorithmic management contexts (Hu et al., 2024), and that acceptance is strongly shaped by perceived control in addition to perceived benefits (Schaap et al., 2024).

These issues are especially salient at first contact, when users have no experiential basis for evaluating an algorithm and must judge based on abstract cues about its decision logic. Stable preferences for decision autonomy can constrain delegation even when delegation is instrumentally attractive (Ivanova-Stenzel and Tolksdorf, 2025), making the perceived loss of choice authority especially salient at the moment of delegation. Task characteristics matter too: trust and reliance depend on perceived objectivity and time pressure (Yang et al., 2025). Additionally, negative reactions to algorithmic errors are robust across different contexts and severity levels (Renier et al., 2021). Together, these factors suggest substantial initial skepticism about delegating to algorithms.

To understand what shapes these first-contact decisions, we must study observable behavior rather than attitudes alone. Behavioral reliance and self-reported trust need not coincide; evidence from AI-advisor interactions shows that determinants of behavioral trust can diverge from self-reported trust (Schreibelmayr and Mara, 2025). Moreover, even in repeated interaction settings, the mere framing of advice as AI-generated can produce systematic overreliance (Klingbeil et al., 2024), underscoring that label-driven cues persist beyond first contact. This motivates our focus on observable delegation behavior, treating first contact as a critical moment when abstract cues about an algorithm’s decision logic are especially influential in shaping whether users delegate.

In this study, we contrast two decision rules that are both plausible in financial choice settings but may elicit different reactions. One algorithm follows

an *objective optimization rule* and selects the option with the highest expected value (EV), which may appear impersonal and easier to justify as a rule-based benchmark. The other is *preference-based personalization* in a narrow sense: it predicts a participant’s preferred choice from their previously elicited preferences rather than maximizing EV. In implementation terms, the personalized algorithm maps each participant’s earlier choices in a simpler elicitation environment to corresponding choices in a more complex environment later in the experiment; it does not involve adaptive updating or learning during the delegation task. The central question is whether people are more willing to delegate to EV maximization than to preference-based personalization. This comparison connects to evidence that responsibility-sharing considerations can shape advisor choice in human-versus-algorithm settings (Gazit et al., 2023) and to recent studies of financial delegation and outcome evaluation (Ismagilova and Ploner, 2025). Related findings in financial delegation are also suggestive of responsibility-sharing considerations (Werner et al., 2026), while evidence on perceived responsibility differences between human and machine partners is mixed (Kirchkamp and Strobel, 2019).

Preregistered hypothesis and analysis plan

Consistent with our preregistration (<https://aspredicted.org/7un5qf.pdf>), the main hypothesis is:

H 1. People are more likely to delegate financial decisions to an EV-maximizing algorithm than to a personalized algorithm that incorporates their (revealed) preferences.

Our preregistered primary analyses test this hypothesis in two complementary ways. First, we test the treatment difference in delegation in the first decision (round 1) using a one-sided Boschloo exact test. This isolates first contact, when participants have no outcome feedback and must decide whether to delegate based on their interpretation of the algorithm. Second, we test whether EV maximization leads to higher delegation overall by comparing participants’ average delegation frequencies across rounds using a one-sided Wilcoxon–Mann–Whitney test.

To account for the repeated-decision structure and additional determinants of delegation, our preregistration also specifies a multivariate regression analysis. We estimate a model with treatment indicators interacted with the decision period, decision-ID fixed effects, and participant-clustered standard errors. We further include the feedback history in the form of including the share of zero outcomes among prior delegated choices, to control for how experienced payoffs may influence subsequent delegation decisions. This preregistered specification

tests whether any treatment difference varies systematically over decisions while controlling for option-set heterogeneity and participants’ feedback histories.

Study overview and supplementary analyses

We implement our design in a between-subjects experiment in which participants repeatedly choose whether to decide themselves or delegate to the assigned algorithm. The repeated-decision setting provides a direct behavioral measure of delegation and allows us to distinguish first-contact delegation from later decisions.

In addition to the preregistered tests, we report a small set of analyses to characterize heterogeneity in delegation propensity. In particular, we examine categorical non-delegation (never delegating across rounds). This outcome captures a distinct form of resistance to algorithmic control transfer and is relevant for understanding which users may opt out of decision aids altogether.

In sum, this paper provides a preregistered test of a simple but consequential claim: even in a setting where preference-based personalization may appear normatively attractive, people may nonetheless be more willing to delegate to a ‘more objective’ algorithm. Our design and preregistered analyses allow us to assess this claim both at first contact and on average across repeated decisions, and show that resistance to personalization can take the form of both first-contact hesitation and categorical non-delegation.

2 Related Work

Our work sits at the intersection of literatures on (i) the behavioral uptake of algorithmic decision aids, (ii) delegation, agency, and responsibility in human–algorithm interaction, (iii) transparency as an interpretive cue for acceptance, (iv) personalization as a design feature of algorithmic decision rules, and (v) heterogeneity in trust and delegation. Below, we situate our study within these streams, identify the specific gap our experiment addresses—the missing comparison between personalized and objective decision rules—and connect them to the robo-advisor application domain.

2.1 Algorithm aversion, appreciation, and conditions for uptake

Behavioral uptake of algorithmic advice is not one-directional: studies document both *algorithm aversion* and *algorithm appreciation*, depending on context, feedback, and comparison benchmark (Dietvorst et al., 2018; Logg et al., 2019). A core

implication is that reliance cannot be predicted based on objective performance alone. Users need to make sense of how recommendations are generated before accepting them, and understandability can gate uptake independently of accuracy (Yeomans et al., 2019). Systematic evidence further shows that uptake is moderated by factors at multiple levels (algorithm, individual, and task), including perceived control and decision context (Mahmud et al., 2022; Chugunova and Sele, 2022).

This mixed-evidence baseline motivates our design choice to study delegation as an observed behavioral outcome and to isolate one concrete determinant of uptake: the algorithmic *decision rule* (objective EV maximization versus preference-based personalization).

2.2 Delegation, agency, and responsibility sharing

Delegation goes beyond advice-taking: it is a transfer-of-control decision that entails ceding choice authority to the algorithm (Candrian and Scherer, 2022). For that reason, agency and responsibility become central concerns. In advisor-choice settings, people evaluate more than expected payoff, including whether responsibility for consequences can be shared with or offloaded onto the advisor (Gazit et al., 2023). Related evidence from financial delegation suggests that delegation patterns are consistent with responsibility-sharing motivations, particularly in loss contexts (Werner et al., 2026). Yet, evidence on whether perceived responsibility differs between human and machine partners is mixed, with Kirchkamp and Strobel (2019) finding little difference in perceived responsibility or guilt.

This control-and-responsibility perspective is consistent with adjacent evidence: perceived trustworthiness dimensions shape acceptance of algorithmic decision-makers (Höddinghaus et al., 2021); preserving human involvement can increase uptake but may lower accuracy (Sele and Chugunova, 2024); and stable preferences for decision autonomy can constrain delegation even when delegation is instrumentally attractive (Ivanova-Stenzel and Tolksdorf, 2025). Taken together, this evidence implies that delegation is shaped by perceived control and responsibility allocation, not only expected returns. This motivates our focus on observed delegation behavior and makes the EV-versus-personalized comparison directly informative, because the two rules differ in how strongly they cue impersonal objectivity versus personal ownership of outcomes.

2.3 Transparency, and interpretive cues for acceptance

Personalization and transparency are often presented as acceptance-enhancing design features, but recent evidence suggests they can also trigger resistance. In recommendation contexts, greater transparency can backfire when users infer

socially negative meanings (for example, feeling judged or dehumanized), reducing acceptance (Chen et al., 2024). In algorithmic management, transparency can likewise show a non-monotonic effect, where additional transparency eventually increases resistance rather than compliance (Hu et al., 2024). Related evidence further indicates that acceptance is driven less by whether the decision-maker is labeled algorithmic or human, and more by expected benefits and perceived control (Schaap et al., 2024).

The broader explainability literature is consistent with this interpretation: transparency and explanations matter, but their effects depend on what they change psychologically. Fairness/accountability/transparency perceptions shape trust and adoption (Shin and Park, 2019); explanations can improve trust calibration and behavior in some settings (Leichtmann et al., 2023); and combining explanation with decision control can improve compliance, though design complexity can also offset those gains (Westphal et al., 2023). In applied decision-support experiments, transparency influences advice use via trusting beliefs and perceived understanding, with effects that depend on the information environment and experienced system performance (Ning et al., 2024; Sun et al., 2026). For our setting, this implies that users’ interpretation of decision logic can shape delegation independently of objective expected returns.

2.4 The missing comparison: personalized versus objective decision rules

Despite the breadth of the literatures reviewed above, a striking gap remains: to our knowledge, no study has experimentally tested whether making an algorithm’s decision rule *preference-based* (as opposed to following an impersonal, objective criterion) changes actual delegation behavior. Several lines of work come close but address different questions.

First, the “algorithm modification” paradigm pioneered by Dietvorst et al. (2018) shows that giving users even slight control over an algorithm’s output increases adoption. This demonstrates that *user input* matters, but the manipulation concerns users’ ability to adjust the output, not whether the underlying decision rule incorporates their preferences or maximizes an objective criterion.

Second, the robo-advisor design literature examines how features such as conversational onboarding (Hildebrand and Bergner, 2021), user controllability over rebalancing (Kim et al., 2025), and the performance–control trade-off in configuration choices (Rühr, 2020) shape adoption and trust. These studies vary the *interface* and the *degree of control*, but hold the nature of the decision rule itself constant or leave it as an implicit design feature.

Third, theoretical and empirical work on robo-advisor personalization has fo-

cused on optimizing portfolio allocation through more granular client interaction (Capponi et al., 2022) or on comparing algorithm-selected portfolios with self-directed choices (Boulou-Reshef et al., 2026). These contributions are informative about *performance* implications of personalization, but do not ask whether users are more or less willing to delegate when the algorithm’s rule is personalized.

This gap is consequential beyond academic interest. In practice, financial robo-advisors, healthcare decision aids, and algorithmic recommender systems routinely face the design choice between objective optimization criteria (such as expected-value or evidence-based guidelines) and preference-based rules that incorporate individual user profiles. If personalization reduces rather than increases the willingness to delegate—as our evidence suggests—then system designers may inadvertently raise adoption barriers by prioritizing preference-matching. Understanding this trade-off is directly relevant for any domain where algorithmic decision aids must be voluntarily adopted to be effective.

2.5 Heterogeneity and contextual moderators of trust and delegation

Recent evidence shows that algorithm uptake is heterogeneous across people, tasks, and decision environments (Ferraz et al., 2025; Yang et al., 2025). At the individual level, delegation varies with psychological readiness and other personal traits linked to delegation behavior (Kim, 2026; Ferraz et al., 2025). Beyond stable traits, uptake also depends on decision architecture: autonomy-preserving choice designs can increase recommendation acceptance (Fink et al., 2024). At the task level, uptake depends on objectivity and time pressure (Yang et al., 2025), while overreliance risks remain salient when users follow advice against their own information (Klingbeil et al., 2024). Taken together, these cross-level moderators imply that average delegation rates can mask meaningful subgroup patterns. For our study, this motivates looking beyond means and explicitly reporting categorical non-delegation as a distinct acceptance pattern.

2.6 Financial advice and robo-advisors as an application domain

Although our task is experimentally stylized, it maps naturally to financial delegation settings where users decide whether to retain control or hand choice authority to an algorithmic decision aid. This is exactly the context in which objective optimization and preference-based personalization often coexist in practice. Existing robo-advisor evidence is informative but fragmented: studies document adoption patterns and how users combine advisor recommendations with their

own judgment (Figà-Talamanca et al., 2022; Mahmud et al., 2025), outcome bias in the evaluation of delegated financial decisions, whether made by a human or an algorithm (Ismagilova and Ploner, 2025), and potential performance gains from robo-advice relative to self-directed choices (Lambrecht et al., 2023).

This makes financial decision-making a strong test bed for our question: whether changing the algorithmic decision rule changes willingness to delegate, even when both rules are plausible.

2.7 Synthesis and contribution

Across these literatures, the common message is that delegation to algorithms is jointly shaped by decision-rule design, perceived control and responsibility, trust calibration, and context-dependent interpretation of transparency and autonomy-related design cues (Dietvorst et al., 2018; Chugunova and Sele, 2022; Schaap et al., 2024; Fink et al., 2024; Hu et al., 2024). Yet as we have documented above, the existing evidence base contains a notable gap: while adjacent work has studied algorithm modification, controllability, conversational design, and performance implications of personalized portfolios, no study has experimentally isolated the effect of the decision rule itself—specifically, whether making the rule preference-based rather than objectively optimizing changes users’ willingness to delegate. This gap matters practically because robo-advisors, healthcare decision aids, and algorithmic recommender systems routinely face exactly this design choice, and a wrong assumption about the direction of the effect could raise rather than lower adoption barriers.

Our study addresses this gap with a preregistered behavioral test in a repeated-decision delegation paradigm, contrasting an EV-maximizing rule with a preference-based rule. By jointly examining first-contact delegation, average delegation across rounds, and categorical non-delegation, we connect classic aversion/ appreciation evidence with newer findings on autonomy, transparency backfire, and responsibility attribution (Logg et al., 2019; Fink et al., 2024; Hu et al., 2024; Gazit et al., 2023).

3 Method

3.1 Participants and preregistered exclusions

Participants were recruited via Prolific from a UK-based participant pool and completed the study online. The preregistered target was 120 participants per treatment condition before exclusions. We recruited $N = 240$ participants and then applied preregistered exclusion rules for data quality and protocol compliance.

Exclusions covered failed bot/slider checks at first attempt, failed comprehension checks after the second attempt, failure to confirm careful participation, disclosed AI usage during the study, implausibly short completion times under the preregistered log-time rule, and incomplete data. The final analysis sample is $N = 233$ (EV condition: $n = 118$; preference-based condition: $n = 115$).

3.2 Design and task

The study used a between-subjects design with two algorithm conditions. In both conditions, participants completed a two-part experiment programmed in the online-ready version of z-Tree (Fischbacher, 2007).

Part 1 consisted of 12 binary lottery choices used to elicit revealed risk preferences. Participants were truthfully told that "[t]he questions in this part are **hypothetical**, but the choices **may be meaningful for Part 2** (recall that your bonus payment from the experiment depends on what happens in Part 2)." The highlighted words were boldfaced in the original instructions. Options were ordered by payment amounts in Part 1 to ease comparison, as can be seen in the upper panel of Figure 1.

In Part 2, participants completed eight incentivized lottery decisions between eight lotteries each. The lottery decisions were connected to the first eight lottery decisions of Part 1 in the following way (the order of tasks was always fully randomized, except that the relevant decisions from Part 1 had to be the first eight decisions out of the 12 decisions in that Part). Figure 1 showcases the relationship between the corresponding tasks. For the Part-2 task in the lower panel, we shuffled payments and states for each Option separately. Then, we added three dominated Options for each of the initial Options, and finally changed the order in which the Options appeared on the screen. We chose not to randomize the orders to prevent tasks from becoming too easy, for example, due to all Options corresponding to one of the two dominating Options appearing without interruption. Even if such an ordering would have been unlikely *ex ante*, it could have made a participant aware of the construction principle of the Part-2 tasks, which is something that we strongly wanted to avoid.

The procedure for each round of Part-2 was as follows: Participants saw the full table for five seconds. Then, they saw a different screen on which they had five seconds to choose between delegating to the algorithm corresponding to their treatment condition and choosing an Option themselves. In the EV condition, delegation selected the lottery with the highest expected value. In the preference-based condition, delegation followed each participant's elicited preference profile from Part 1. Assignment to condition was randomized at the participant level. In the EV condition, participants were informed that "[the] algorithm works out how much each option would yield on average, and selects the option for which

	State 1	State 2	State 3	State 4	State 5	State 6	State 7	State 8	State 9	State 10
Option 1	225	225	225	225	225	225	225	205	0	0
Option 2	525	525	525	525	525	0	0	0	0	0

State:	1	2	3	4	5	6	7	8	9	10
Option 1	225	225	0	225	0	225	225	225	225	0
Option 2	525	0	505	0	525	0	0	0	525	505
Option 3	525	0	525	0	525	0	0	0	525	525
Option 4	225	225	0	225	205	205	205	225	225	0
Option 5	225	205	0	205	205	225	225	205	225	0
Option 6	225	225	0	225	205	225	225	225	225	0
Option 7	525	0	525	0	525	0	0	0	225	525
Option 8	505	0	525	0	505	0	0	0	505	525

Figure 1: Example screens from Part 1 (preference elicitation, upper panel) and Part 2 (delegation decisions with lottery choices, lower panel).

this is highest.” In the preference-based condition, they instead read that “[the] algorithm analyses your choices from Part 1 and selects the best option for you in light of your Part-1 choices.”

If they did not delegate, they saw the table containing the eight Options for up to eight more seconds, during which they had to choose one of the Options. If they did not make a choice in the payoff-relevant choice, they did not earn a bonus payment from the experiment. The total payoff from the experiment was GBP 2.00 plus the points earned in one randomly selected round of Part 2, converted at a rate of 1 penny per point (participants could earn up to GBP 8.90 in addition to the base payoff). On average, participants earned GBP 4.14 including the base payoff for a median time of just under 16 minutes.

3.3 Procedure

After consent, participants completed attention and comprehension checks, then Part 1 (preference elicitation), followed by a question on how satisfied they were with their choices from Part 1. Then, participants completed Part 2 (delegation decisions with round-by-round outcome feedback), followed by a question of whether they would delegate to the algorithm in case they had to choose to delegate or not for 10 (hypothetical) additional tasks. After indicating that choice, participants estimated (on a 0–10 integer scale) in how many of those 10 tasks the algorithm would choose a non-preferred option and in how many it would choose a “completely weird” option. Before the final post-task questionnaire

containing demographics and attitudinal items, we asked them for self-reported reasons for their choice in the hypothetical choice at the end of Part 2. Because our theoretical focus is reliance behavior, primary outcomes are behavioral rather than attitudinal.

3.4 Ethics approval

The study received ethics approval from the Institutional Review Board of the Gesellschaft für experimentelle Wirtschaftsforschung e.V. (GfEW; German Association for Experimental Economic Research), certificate no. 19HBJNwQ (status: approved; expedited review; approval date: 11/13/2025; certificate validity through 11/13/2027; see <https://gfew.de/ethik/19HBJNwQ>). All participants provided informed consent before participation.

3.5 Preregistered outcomes and analyses

The preregistered hypothesis (H1) predicted higher delegation to the EV-maximizing algorithm than to the preference-based algorithm. The preregistration document is available at <https://aspredicted.org/7un5qf.pdf>.

Primary preregistered analyses were:

1. a one-sided Boschloo exact test for the treatment difference in delegation in round 1 (first contact), and
2. a one-sided Wilcoxon–Mann–Whitney test comparing participant-level average delegation rates across Part 2.

We additionally estimated a preregistered linear probability model with decision-set fixed effects and participant-clustered standard errors, including treatment-by-period interactions and the share of outcomes with a payoff of zero among prior delegated choices. We further estimated an exploratory specification excluding round 1 to separate first-contact effects from choices made after feedback experience.

4 Results

4.1 Overview

Two empirical patterns organize the results. First, treatment differences are concentrated at first contact. Second, categorical non-delegation is substantially more common under preference-based personalization.




Figure 2: Average delegation over decisions in Part 2 by treatment condition.

Choice quality among non-delegated decisions was low in both treatments. The share of non-dominated self-chosen options was 34.1% in the EV condition and 35.1% in the preference-based condition, implying that roughly two-thirds of self-made choices were dominated. Hence, participants had a clear instrumental reason to delegate. On average, delegated decisions yielded 243.49 points (GBP 2.43), compared to 195.18 points (GBP 1.95) for non-delegated decisions, a difference of approximately 25%.

4.2 Primary test: first-contact delegation

In round 1, delegation was higher in the EV condition than in the preference-based condition (66.9% vs. 55.8%). The preregistered one-sided Boschloo exact test yields $p = 0.0413$, supporting H1 for first-contact behavior. Figure 2 visualizes the round-level profile and shows that the treatment gap is concentrated in the first decision.

4.3 Primary test: delegation across repeated decisions

At the participant level, the preregistered one-sided Wilcoxon–Mann–Whitney test on average delegation across Part 2 yields $W = 6428$ and $p = 0.2425$. We therefore do not find evidence of a treatment difference in average delegation across rounds.

The preregistered regression specification points in the same direction. The EVAlgorithm coefficient is small and statistically insignificant ($\beta = -0.0487$, SE




Figure 3: ECDF of participant-level average delegation by treatment condition.

= 0.0831, $p = 0.558$). Likewise, both interaction terms are imprecisely estimated (EValgorithm $\times$ WithinPartDecisionNumber: $\beta = 0.0019$, SE = 0.0115, $p = 0.871$; EValgorithm $\times$ share[Payoff = 0]: $\beta = -0.0497$, SE = 0.0977, $p = 0.612$), indicating no robust treatment difference over repeated decisions. By contrast, the share of prior delegated outcomes with a payoff of zero is strongly negatively associated with subsequent delegation ($\beta = -0.3854$, SE = 0.0781, $p < 0.001$).³

Taken together, the preregistered tests indicate that the treatment effect is concentrated at first contact and does not persist over repeated decisions, while negative feedback (in particular prior zero-outcomes) is strongly associated with lower subsequent delegation. The convergence in delegation behavior over rounds is visible in Figure 2; the cross-participant distribution is shown in Figure 3.

4.4 Categorical non-delegation

This section reports an outcome that was not preregistered (categorical non-delegation) that is complementary to the preregistered primary tests.

The second clear pattern is categorical non-delegation, which can be seen from the left end of Figure 3. In the preference-based condition, 18 of 115 participants never delegated (15.7%); in the EV condition, this applies to only 4 of 118 participants (3.4%). This gap points to a distinct acceptance barrier under personalization for a subset of users.⁴ Note that although there also seems to be a

³Running an analogous regression with fixed effects for the task on round-1 data only yields the same result as the Boschloo test, with $\beta = 0.113$, SE = 0.065, and $p = 0.041$ (one-sided).

⁴The same direction appears in an earlier study with a different design (reported in Suraci’s PhD

treatment difference in always-delegators, this difference is far less pronounced (19.1% vs. 13.6%).

Never-delegators were not meaningfully different in age from participants who delegated at least once (mean age 41.8 vs. 42.9 years; $t = -0.329$, $p = 0.745$), and we likewise find no clear association with household income (categorical; $\chi^2 = 13.171$, $df = 11$, $p = 0.282$) or education (categorical; $\chi^2 = 1.859$, $df = 5$, $p = 0.868$), and both spent similar amounts of time in the experiment (mean number of seconds 1113 for delegators vs. 1237 for never-delegators, $t = -0.971$, $p = 0.341$). And while there is a tendency for male participants to be never-delegators more often than females (12.0% vs. 6.3%), the evidence is weak (Boschloo exact test, $p = 0.156$). We also do not find meaningful differences in self-reported willingness to take financial risks (4.53 vs. 4.68; $t = -0.224$, $p = 0.825$), self-reported response reliability/care (Reliable: 6.91 vs. 6.91; $t = -0.059$, $p = 0.954$), or frequency of AI use/interactions (AIXperience, categorical: 5.87 vs. 5.18; $t = 0.866$, $p = 0.395$). Within the preference-based condition, however, never-delegators reported higher satisfaction with their Part-1 choices than delegators (7.94 vs. 7.56; $t = -1.400$, $p = 0.169$), which is inconsistent with the interpretation that non-delegation mainly reflects dissatisfaction with the elicited preference profile used by the personalized algorithm.

4.5 Responses to feedback

This preregistered secondary analysis characterizes behavioral dynamics rather than serving as an additional primary test of H1.

To characterize feedback dynamics, we estimate a win-stay/lose-shift model on rounds 2–10. Under a preference-based algorithm, a prior win (a payoff larger than zero) is associated with a higher—though imprecisely estimated—probability of staying with the same delegation/non-delegation choice ($\beta_{\text{lagWin}} = 0.043$, $SE = 0.033$, $p = 0.194$). The interaction between the EV treatment and a prior win is negative and significant ($\beta = -0.102$, $SE = 0.049$, $p = 0.041$), implying that in the EV condition, participants tend to shift their delegation/non-delegation decision after a non-0 payoff *more* often than after a zero-payoff (approximately $0.043 - 0.102 = -0.059$; still, even in this case, stay rates are 60%; also note that average stay rates after losses are 69% in the Preference-based condition and 65% in the EV condition, so that we generally do not observe a win-stay–lose-shift pattern).

This dynamics result reinforces the broader pattern: treatment differences are phase-specific and subgroup-specific rather than visible in aggregate average delegation.

thesis): 32/189 never-delegators in preference-based treatments versus 17/187 in EV treatments (one-sided Boschloo exact test, $p = 0.0135$).




Figure 4: Reasons for hypothetical delegation choices by treatment and delegation type. Note: The exact questions were, after choosing (not) to delegate, *BetterJob*: I thought the computer could do a better job than me (I could do a better job than the computer); *DislikeRelyingOn*: I don't like relying on myself (computers); *DontTrust*: I don't fully trust myself (the computer); *HowExactly*: I find it unpleasant that I don't know how exactly the computer chooses an option for me; *MoreInfo*: I would have preferred more information about how the computer makes its decisions; *Other*: Other (please specify below); *PerformanceCuriosity*: I wanted to see how well the computer (I) can do; *Reactance2Analysis*: I don't like to be analysed by an algorithm; *Responsibility4Choice*: I didn't want (wanted) to be responsible for choosing an (the) option; *Responsibility4Outcome*: I didn't want to be responsible if a low payment results; *Tedious*: Doing it myself would have been too tedious.

4.6 Exploratory mechanisms: post-task reasons and interpretations

To better understand the behavioral patterns above, we analyze participants' self-reported reasons collected after Part 2. Specifically, after the incentivized rounds, participants indicated whether they would delegate in 10 hypothetical additional tasks and then reported reasons for that hypothetical choice (closed-item checklist plus open text). We treat these responses as exploratory mechanism evidence. Figure 4 shows reason-selection frequencies by treatment and delegation type.

The first thing to note is that reason-selection patterns vary more by delegation type than by treatment. The only visible treatment contrast is that differences between never-delegators and delegators appear larger in the EV condition than in the preference-based condition. We do not interpret this contrast because the

never-delegator subgroup in the EV condition is very small (4 participants).

Two descriptive patterns are consistent with the behavioral results. First, participants who delegated (at least once) were more likely to indicate performance-based motives (on average, 70.1% indicated that the computer would choose better, while only 36.4% of the never-delegators indicated this). Second, participants who never delegated were more likely to select reasons centered on distrust of the algorithm (32.2% vs. 50%) and retaining personal responsibility for the choice (38.9% vs. 63.6%). This mirrors the treatment-level non-delegator pattern and is consistent with the idea that a subgroup places relatively more weight on agency and responsibility concerns than on expected performance.

Exploratory interaction models lend quantitative support to treatment-specific deterrence channels (Figure 5 in the appendix). In a linear probability model of individual delegation decisions interacting treatment with participants' algorithmic performance expectations (controlling for decision round, per-choice Part-1 confidence, revealed Part-1 risk preferences, and task fixed effects), the non-preferred-choices estimate is strongly negatively associated with delegation in the preference-based condition ($\beta = -0.058$, $SE = 0.016$, $p < 0.001$) and substantially attenuated in the EV condition (EValgorithm $\times$ non-preferred: $\beta = +0.035$, $SE = 0.020$, $p = 0.090$). Conversely, the completely-weird-choices estimate shows near-zero association with delegation in the preference-based condition ($\beta = +0.007$, $SE = 0.018$, $p = 0.709$) but is strongly and negatively associated with delegation in the EV condition (EValgorithm $\times$ completely weird: $\beta = -0.070$, $SE = 0.023$, $p = 0.003$). This pattern suggests that what deters delegation differs by algorithm type: perceived non-preferred choices are a deterrent specifically under personalization, while perceived extreme errors (completely weird choices) are a deterrent under objective optimization.

Additional reason items are informative by contrast. Items such as not wanting to be analyzed by an algorithm were rarely selected overall, suggesting that profiling concerns or dehumanization were unlikely to be primary drivers of treatment differences in this setting. By comparison, selecting the task being tedious appears more among delegators (though still at a low level, 10.9% vs. 0%), consistent with an effort-saving channel among participants who are willing to hand over control. Open-text responses also repeatedly mention time pressure and limited time to scan all options; these comments point to a general choice-friction channel, but they do not isolate a treatment-specific mechanism.

These mechanism patterns should be interpreted cautiously. The reason measures are ex-post and refer to hypothetical additional choices rather than the incentivized rounds themselves. They therefore function as suggestive proxies for motives, not as causal identification of psychological channels or as decisive evidence for one mechanism over another.

5 Discussion

5.1 Main takeaway

This preregistered study shows that preference-based personalization can deter delegation rather than increase it. The deterrence appears in two forms: lower delegation at first contact and a larger subgroup of participants who never delegate under personalization.

The first-contact finding is important because it captures behavior before participants can learn from outcome feedback. At that stage, delegation is driven primarily by how participants interpret the algorithmic rule itself. Our evidence is consistent with the interpretation that an impersonal EV rule is initially easier to accept as a legitimate external computation than a profile-based personalized rule based on the user’s own revealed preferences.

5.2 Interpretation in relation to prior work

Our main contribution is showing that algorithm uptake depends on the decision rule, not just on whether a system is algorithmic. The contrast between expected-value maximization and preference-based personalization reveals that users may resist personalization despite its normative appeal.

The categorical non-delegation pattern is particularly informative: many participants entirely avoid delegation under personalization, consistent with work on autonomy, responsibility, and control as drivers of acceptance. Exploratory post-task reasons (Section 4.6) align with this interpretation, especially the stronger emphasis among never-delegators on responsibility and distrust, while concerns about being algorithmically analyzed appear comparatively rare. Because these reason measures are ex-post and hypothetical, this evidence remains suggestive rather than causal.

These findings are consistent with an earlier study on the same question (reported as a chapter in Suraci’s PhD thesis), which motivated and informed the design of the present preregistered experiment. In that study, the delegation gap between EV and preference-based treatments was larger (about 20 percentage points) and persisted over decisions. Recomputing the categorical non-delegation contrast for that earlier sample also yields a significant difference (32/189 in preference-based treatments vs. 17/187 in EV treatments; one-sided Boschloo exact test, $p = 0.0135$).⁵

⁵Note that categorical non-delegation was not analyzed or reported in the earlier study, which is the reason for why we did not pre-register its analysis for the present study.

5.3 Limitations

Several limitations should be noted. First, the study is an online experiment with stylized financial lotteries, so external validity to field settings should be interpreted cautiously. Second, although behavioral outcomes are well suited for studying reliance, the design does not causally identify the psychological mechanism behind first-contact hesitation or categorical non-delegation (e.g., perceived objectivity, profiling concerns, or accountability salience). Third, the present analyses focus on short-run repeated interaction; longer-run adaptation to personalized systems remains an open question.

5.4 Implications and future research

For design, the results imply that personalization should not be treated as a uniformly beneficial acceptance lever. In settings where objective criteria are salient and normatively legitimate, objective optimization may face lower initial resistance than preference-based personalization. This has direct implications for the design and communication of decision aids in finance and other high-stakes domains.

These implications extend beyond our experimental setting. Financial robo-advisors routinely face the choice between objective portfolio rules (e.g., mean-variance optimization) and preference-based personalization derived from risk profiling questionnaires (Capponi et al., 2022; Rühr, 2020). Our findings suggest that the very feature intended to increase subjective fit—tailoring the algorithm to the user’s profile—may inadvertently encroach on users’ sense of choice authority, deterring precisely those users who value exercising active control over their decisions. The same tension arises whenever algorithmic decision aids must be voluntarily adopted: in healthcare, treatment recommender systems can follow clinical-evidence guidelines or patient-preference protocols; in hiring, screening tools can apply standardized criteria or adaptively weight candidate profiles. In each case, the assumption that preference alignment promotes adoption deserves empirical scrutiny rather than being taken as a design axiom.

Future research should test mechanism-targeted interventions at first contact, including reframing of personalization logic, responsibility framing, and alternative explanation formats. It should also examine whether the non-delegator subgroup is stable across domains, stakes, and cultural contexts.

5.5 Contributions

This paper contributes in three ways. First, it provides preregistered evidence that personalization can reduce, rather than increase, delegation at first contact.

Second, it shows that resistance is not only transient: personalization is associated with substantially more categorical non-delegation. Third, by comparing two algorithmic decision rules within the same task, it clarifies that behavioral uptake depends on perceived rule characteristics, not only on whether advice is algorithmic or human.

5.6 Data and code availability

Data, analysis code, and study materials will be deposited in a public repository upon acceptance/publication of this manuscript.

Appendix




Figure 5: Condition-specific marginal effects on delegation probability. The plot shows all regression coefficients from the linear probability model (main effects and treatment interactions). Key expectation variables include the number of non-preferred choices (the participant’s own prediction of how many tasks, out of 10, the algorithm would choose differently from their preference) and the number of completely weird choices (predicted count of highly unusual choices). The preference-based-condition estimate is the main-effect coefficient; the EV-condition estimate is the sum of the main effect and the treatment interaction term. Error bars are 95% confidence intervals based on participant-clustered standard errors. The model controls for decision round, per-choice Part-1 confidence, revealed Part-1 risk preferences (measured by the number of choices that were less risky in the sense of having fewer states with a payoff of 0), and task fixed effects.

References

- Aysolmaz, B., Müller, R., and Meacham, D. (2023). The public perceptions of algorithmic decision-making systems: Results from a large-scale survey. *Telematics and Informatics*, 79:101954.
- Boulou-Reshef, B., Direr, A., and von Wilczur, N. (2026). Algorithmic vs. human portfolio choice. *Finance*. Forthcoming.

- Candrian, C. and Scherer, A. (2022). Rise of the machines: Delegating decisions to autonomous ai. *Computers in Human Behavior*, 134:107308.
- Capponi, A., Olafsson, S., and Zariphopoulou, T. (2022). Personalized robo-advising: Enhancing investment through client interaction. *Management Science*, 68(4):2485–2512.
- Chen, G., Dang, J., and Liu, L. (2024). After opening the black box: Meta-dehumanization matters in algorithm recommendation aversion. *Computers in Human Behavior*, 161:108411.
- Chugunova, M. and Sele, D. (2022). We and it: An interdisciplinary review of the experimental evidence on how humans interact with machines. *Journal of Behavioral and Experimental Economics*, 99:101897.
- Dietvorst, B. J., Simmons, J. P., and Massey, C. (2018). Overcoming algorithm aversion: People will use imperfect algorithms if they can (even slightly) modify them. *Management Science*, 64(3):1155–1170.
- Ferraz, V., Houf, L., Pitz, T., Schwieren, C., and Sickmann, J. (2025). Trust in the machine: How contextual factors and personality traits shape algorithm aversion and collaboration. *Computers in Human Behavior Reports*, 17:100578.
- Figà-Talamanca, G., Tanzi, P. M., and D’Urzo, E. (2022). Robo-advisor acceptance: Do gender and generation matter? *PLOS ONE*, 17(6):1–13.
- Fink, L., Newman, L., and Haran, U. (2024). Let me decide: Increasing user autonomy increases recommendation acceptance. *Computers in Human Behavior*, 156:108244.
- Fischbacher, U. (2007). z-tree: Zurich toolbox for ready-made economic experiments. *Experimental Economics*, 10(2):171–178.
- Gazit, L., Arazy, O., and Hertz, U. (2023). Choosing between human and algorithmic advisors: The role of responsibility sharing. *Computers in Human Behavior: Artificial Humans*, 1(2):100009.
- Hildebrand, C. and Bergner, A. (2021). Conversational robo advisors as surrogates of trust: Onboarding experience, firm perception, and consumer financial decision making. *Journal of the Academy of Marketing Science*, 49(4):659–676.
- Hu, P., Zeng, Y., Wang, D., and Teng, H. (2024). Too much light blinds: The transparency-resistance paradox in algorithmic management. *Computers in Human Behavior*, 161:108403.

- Höddinghaus, M., Sondern, D., and Hertel, G. (2021). The automation of leadership functions: Would people trust decision algorithms? *Computers in Human Behavior*, 116:106635.
- Ismagilova, Z. and Ploner, M. (2025). Ain't blaming you: Delegation of financial decisions to humans and algorithms. *Computers in Human Behavior: Artificial Humans*, 4:100147.
- Ivanova-Stenzel, R. and Tolksdorf, M. (2025). Delegating in the age of ai: Preferences for decision autonomy. Discussion Paper 558, Ludwig-Maximilians-Universität München und Humboldt-Universität zu Berlin, Collaborative Research Center Transregio 190 - Rationality and Competition, München und Berlin.
- Kim, H. (2026). Beyond algorithm aversion: The impact of psychological readiness on algorithmic advice. *Computers in Human Behavior*, 174:108824.
- Kim, S., Ahn, D., Choi, B., Park, J., Yang, J., Yoo, B., and Yum, K. (2025). Controllability in algorithmic decision-making: Field evidence from a robo-advisor. In *Proceedings of the International Conference on Information Systems (ICIS 2025)*. AIS.
- Kirchkamp, O. and Strobel, C. (2019). Sharing responsibility with a machine. *Journal of Behavioral and Experimental Economics*, 80:25–33.
- Klingbeil, A., Grützner, C., and Schreck, P. (2024). Trust and reliance on ai — an experimental study on the extent and costs of overreliance on ai. *Computers in Human Behavior*, 160:108352.
- Lambrecht, M., Oechssler, J., and Weidenholzer, S. (2023). On the benefits of robo advice in financial markets. AWI Discussion Paper Series 734, Heidelberg University.
- Leichtmann, B., Humer, C., Hinterreiter, A., Streit, M., and Mara, M. (2023). Effects of explainable artificial intelligence on trust and human behavior in a high-risk decision task. *Computers in Human Behavior*, 139:107539.
- Logg, J. M., Minson, J. A., and Moore, D. A. (2019). Algorithm appreciation: People prefer algorithmic to human judgment. *Organizational Behavior and Human Decision Processes*, 151:90–103.
- Mahmud, H., Islam, A. N., Ahmed, S. I., and Smolander, K. (2022). What influences algorithmic decision-making? a systematic literature review on algorithm aversion. *Technological Forecasting and Social Change*, 175:121390.

- Mahmud, H., Islam, N., and Krishnan, S. (2025). Human-robo-advisor collaboration in decision-making: Evidence from a multiphase mixed methods experimental study. *Decision Support Systems*, 198:114541.
- Ning, X., Lu, Y., Li, W., and Gupta, S. (2024). How transparency affects algorithmic advice utilization: The mediating roles of trusting beliefs. *Decision Support Systems*, 183:114273.
- Renier, L. A., Schmid Mast, M., and Bekbergenova, A. (2021). To err is human, not algorithmic – robust reactions to erring algorithms. *Computers in Human Behavior*, 124:106879.
- Rühr, A. (2020). Robo-advisor configuration: An investigation of user preferences and the performance-control dilemma. In *Proceedings of the 28th European Conference on Information Systems (ECIS 2020)*. AIS.
- Schaap, G., Bosse, T., and Hendriks Vettehen, P. (2024). The abc of algorithmic aversion: not agent, but benefits and control determine the acceptance of automated decision-making. *AI & SOCIETY*, 39(4):1947–1960.
- Schreibelmayr, S. and Mara, M. (2025). Determinants of self-reported and behavioral trust in an ai advisor within a cooperative problem-solving game. *Computers in Human Behavior: Artificial Humans*, 6:100235.
- Sele, D. and Chugunova, M. (2024). Putting a human in the loop: Increasing uptake, but decreasing accuracy of automated decision-making. *PLOS ONE*, 19(2):1–14.
- Shin, D. and Park, Y. J. (2019). Role of fairness, accountability, and transparency in algorithmic affordance. *Computers in Human Behavior*, 98:277–284.
- Sun, Y., Liao, M. M., Sundar, S. S., and Walther, J. B. (2026). Does transparency matter when an ai system meets performance expectations? an experiment with an online dating site. *Computers in Human Behavior*, 177:108875.
- Werner, J., Suraci, L., and Wolff, I. (2026). In losses we trust: Delegating decisions to artificial intelligence. Mimeo.
- Westphal, M., Vössing, M., Satzger, G., Yom-Tov, G. B., and Rafaeli, A. (2023). Decision control and explanations in human-ai collaboration: Improving user perceptions and compliance. *Computers in Human Behavior*, 144:107714.
- Yang, R., Li, S., Qi, Y., Liu, J., He, Q., and Zhao, H. (2025). Unveiling users’ algorithm trust: The role of task objectivity, time pressure, and cognitive load. *Computers in Human Behavior Reports*, 18:100667.

Yeomans, M., Shah, A., Mullainathan, S., and Kleinberg, J. (2019). Making sense of recommendations. *Journal of Behavioral Decision Making*, 32(4):403–414.