To Reviewer 3 (report not organized by the 9 standardized Questions):

First of all, thank you very much for the thoughtful and fair review. We did our best to address your concerns in the revised version. Below, we reproduce your comments (in black) and provide our answers in blue.

This manuscript under review explores the impact of preference-driven personalized algorithms and objective expectation maximization algorithms on users' algorithmic delegation behavior. It proposes that personalized algorithms may actually reduce users' willingness to delegate, and distinguishes between accuracy-based trust and authority-based resistance paths by combining behavioral data, participants' post-expectation assessments, and self-reported motivations. The research is comprehensive and theoretically sound. The following issues need to be addressed:

Comments

1. The study used a virtual lottery task, not real investment decisions. Does this context affect the results?

We have expanded the "Limitations" subsection to clarify that the Part-2 choices were real and incentivized, not hypothetical vignettes: one round was paid out for real money, contingent on the participant's own choice or, if delegated, on the algorithm's choice. We distinguish this from the task domain, which remains stylized: real investment or medical decisions add ambiguous/delayed payoffs, product complexity, and institutional framing (a named bank or physician) that could independently affect how the loss of choice authority is experienced. We note explicitly that this abstraction is what let us isolate the construal mechanism cleanly, while flagging generalization to richer real-world domains as an open question, cross-referenced to "Future research."

2. The interaction time in this experiment was limited. Can the results be generalized to long-term human-computer interaction in real investment scenarios?

We extended the "Limitations" subsection to acknowledge the within-session narrowing of the treatment gap as suggestive of short-run adaptation, while explicitly refusing to extrapolate to long-term use (months/years vs. minutes). In order not to leave your question completely unanswered, we additionally lay out two competing forces that longer-run field exposure could trigger: (a) accumulating track record and familiarity eroding authority-based resistance in favor of accuracy-based trust, vs. (b) deepening profiling, higher stakes, and switching costs sustaining or amplifying the choice-authority effect. We make clear that our design cannot adjudicate between those, calling for longitudinal field studies as the necessary next step.

3. Figure 2 shows a significant decrease in delegation rates from the first to the second round under EV conditions. The authors' explanation for this phenomenon is biased towards post-expectation interpretation. Does this mean that the authority effect is overridden by performance feedback after the initial contact?

Thank you very much for pressing us on this interpretation, as our resulting re-analysis led to further insights. We do not read the round-1-to-round-2 decline as evidence that feedback overrides the authority effect. We have, however, revised our explanation of it. The submitted version attributed it to the EV rule mechanically producing more zero outcomes. A re-analysis showed that that cannot be the driving force: the reaction to a zero payoff is the same in both conditions (0.254 vs. 0.250), and realized zero shares among delegated decisions differ much less than we thought (40.3% vs. 36.2%, $p = 0.178$). The text now reports, as exploratory analyses, that (i) average

delegation in rounds 2 to 8 seems to be the net effect of two offsetting differences of six to seven points each (more categorical opt-out under personalization, lower delegation under EV among everyone else), and (ii) EV picks that differ from the participant's own Part-1 preference pay zero more often and are followed by less delegation, which accounts for part of the larger EV decline (about six points in round 2, bootstrap 95% CI [0.4, 11.8]). The larger EV decline itself is imprecisely estimated (p = 0.141). Having described the above new findings, the text grants that a substantial amount of the larger decline under EV remains unexplained.

4. The hypothesis regarding the heterogeneity of those who never commission is not fully validated; experimental results show no significant differences between the non-commissioning and commissioning groups on multiple indicators.

Point taken. We are running into power issues here: there are only 22 never-delegators in total (18 preference-based, 4 EV). We have added an explicit caveat in the text noting that these subgroup comparisons are underpowered to detect anything but huge characteristic differences, so the absence of significant contrasts should not be read as evidence that never-delegators are indistinguishable from delegators on these dimensions. We deliberately did not report alternative tests with partially different conclusions (e.g., a KS test on time-in-experiment) to avoid the appearance of test-shopping; the honest characterization is that we are hopelessly underpowered, rather than having evidence of zero-differences.

5. Section 3.5 mentions supplementary analysis, but does not clearly define which analyses are confirmatory and which are exploratory.

Thank you for pointing that out, we have made an extra effort to make it 100% clear what is confirmatory and that is exploratory. To this end, we have restructured the "Responses to feedback" subsection so that the status of every analysis is stated explicitly. In particular, we corrected an imprecision in our earlier version: the text ordering had suggested that the win-stay/lose-shift model described first was the preregistered analysis, which was not true. The preregistered secondary hypothesis was a different one: after delegating to the algorithm with a payoff of zero, participants in the preference-based condition are more likely to switch to non-delegation in the next round than under any other circumstances, in particular than under the same circumstances in the EV condition. The subsection now reports this preregistered test first (it is not supported: the interaction is slightly negative and insignificant, one-sided $p = 0.633$), and then presents the win-stay/lose-shift model and its decomposition by the source of the prior outcome (delegated vs. self-chosen) together, clearly marked as exploratory analyses. The pooled win-stay/lose-shift model itself now sits in a short appendix (Appendix B), with its main results summarized in two sentences in the text. The Introduction, Method, and Discussion sections were adjusted accordingly (the separate "Contributions" subsection was removed in the consolidation requested by Reviewer 1).

6. The author need to carefully review the entire text; for example, there is an error in citing figures such as "Figure?" on page 21.

Good catch, thank you. We decided to take out that figure «last moment» before submitting, and overlooked that pointer to the figure. We are sorry for not checking thoroughly enough. The underlying figure (`ExpectedBadAndAbsurdChoicesByTrmt.pdf`) still exists in the repository if you would prefer we put it back in as an actual figure instead of describing the result in prose only.

7. Section 3.1 lists six pre-registration exclusion criteria, but does not explain the sample distribution corresponding to each exclusion reason when the sample size is reduced from 240 to 233.

We reconstructed the exact breakdown from the raw session data. Of the 7 excluded participants: 1 failed the bot/slider check, 4 failed the careful-participation check (`Reliable < 6`; one of these also disclosed AI usage), and 3 disclosed AI usage during the study (one overlapping with the careful-participation failure). No participant was excluded for a comprehension-check failure or under the log-completion-time rule. We added this breakdown to the Method section.

8. Authors are recommended to cite the following references to enrich the research background.
(i) https://doi.org/10.1186/s40854-026-00941-w
(ii) https://doi.org/10.1186/s40854-024-00671-x

Thank you for these suggestions. We checked both suggested papers and chose to include one of them, Kou et al. (2026). The paper is clearly relevant, as it frames human–AI interaction in finance around the decision architecture and the allocation of authority, which connects naturally to our choice-authority account. We added it as a citation in "Synthesis and contribution" and in the implications discussion. Yüksel et al. (2025) is also relevant, but much further from our topic (no connection to delegation, algorithms, or choice authority). We therefore chose not to cite it, consistent with the Associate Editor's note that suggested references need not all be incorporated.

---

9. Could the manuscript benefit from language editing?

Yes

Sorry to hear that you were not happy with the writing. Beyond the AI-assisted language support already disclosed in the "Declaration of generative AI" subsection, we ran an additional proofreading pass before resubmission by a native speaker.